

Problem Set 1: Applied Stats II

Caitlín Cooney

Due: February 19, 2023

Question 1

Using R generate 1,000 Cauchy random variables (`rcauchy(1000, location = 0, scale = 1)`) and perform the test (remember, use the same seed, something like `set.seed(123)`, whenever you're generating your own data).

```
1 set.seed(123) data <- (rcauchy(1000, location = 0, scale = 1))
```

Write an R function that implements the Kolmogorov-Smirnov test where the reference distribution is normal.

```
1 kolsmir.test <- function(data) {  
2  
3   # Create empirical distribution of observed data  
4   ECDF <- ecdf(data)  
5   empiricalCDF <- ECDF(data)  
6  
7   # Generate test statistic  
8   D <- max(abs(empiricalCDF - pnorm(data)))  
9  
10  # Calculate p-value  
11  p.value <- 1 - pnorm(sqrt(length(data)) * D)  
12  
13  return(list(D = D, p.value = p.value))  
14 }
```

State your null and alternate hypotheses:

H0: The data comes from the specified distribution.

H1: At least one value does not match the specified distribution.

```
1 # Run the function  
2 kolsmir.test(data)
```

The results are:

```
$D
```

```
[1] 0.1347281
```

```
$p.value
```

```
[1] 1.019963e-0}
```

As D is not greater than the critical value, we fail to reject the null hypothesis that the data comes from the specified distribution (or in other words, that $P = P_0$).

Question 2

Estimate an OLS regression in R that uses the Newton-Raphson algorithm (specifically BFGS, which is a quasi-Newton method).

```
1 set.seed (123)
2 data <- data.frame(x = runif(200, 1, 10))
3 data$y <- 0 + 2.75*data$x + rnorm(200, 0, 1.5)

1 linear.lik <- function(theta, y, X) {
2   n <- nrow(X)
3   k <- ncol(X)
4   beta <- theta [1:k]
5   sigma2 <- theta [k+1] ^2
6   e <- y-X%%beta
7   logl <- -.5*n*log( 2*pi )-.5*n*log(sigma2)-((t(e)%*%e)/(2*sigma2))
8   return(-logl)
9 }
10
```

Find parameters that specify the point.

```
1 linear.MLE <- optim(fn = linear.lik, par = c(1, 1, 1),
2                     hessian = TRUE, y = data$y, X = cbind(1, data$x),
3                     method = "BFGS")
4
5 linear.MLE$par
```

The results show:

```
[1] 0.1398324 2.7265559 -1.4390716
```

Show that you get the equivalent results to using `lm`.

```
1 summary(lm(y~x, data))
```

The results show:

```
Coefficients:
Estimate Std. Error t value Pr(>|t|)
(Intercept)  0.13919    0.25276   0.551    0.582
x            2.72670    0.04159  65.564 <2e-16 ***
```

As expected, the estimates for `lm` are the same as the parameters found using the Newton-Raphson algorithm, because ordinary least squares is equivalent to maximum likelihood for a linear model, so it makes sense that `lm` would give us the same answers.